RUH.AI

V1 Design Partner Pilot Blueprint & Rollout Runbook

A repeatable deployment playbook for proving digital employees with real customers, using packs, tenant overlays, approvals, and measured phase gates.

Applies to: Construction Project Operations + AI BuildOps reference packs
Audience: Product, engineering, implementation, customer success, design partners
Status: Working v1 execution document

What this document answers

- Which customers qualify for a design-partner pilot and what must be true before kickoff.
- How a pilot moves from monitor mode to draft mode, approval mode, and constrained autonomy.
- Which metrics, roles, artifacts, and gate reviews determine pilot success and production expansion.

1. Purpose of this document

The previous documents define the platform. This one defines the path to prove it with real customers.

This blueprint converts the Ruh.ai product stack into an executable customer deployment model. It assumes the platform PRD, system design, API/event contracts, employee pack specification, reference pack definitions, and tenant overlay kits already exist. The pilot blueprint is the missing execution layer between product design and production rollout.

Its job is simple: make sure every design partner is qualified consistently, deployed safely, measured against a baseline, and either promoted to production or exited with clean learnings that improve the packs, overlays, and platform.

Qualify	Deploy	Prove
Confirm sponsor, workflow owner, systems of record, baseline metrics, and the trust boundary before work begins.	Compile the pack + tenant overlay, connect systems, run shadow mode, and introduce approval-gated write-backs in phases.	Use a single scorecard for business value, trust, technical health, and adoption before expanding or commercializing.

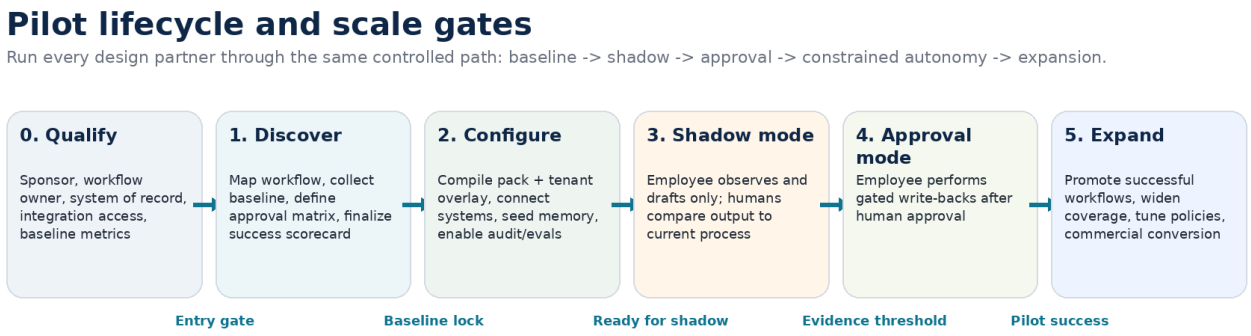

Figure 1. Every pilot follows the same control path and only expands after evidence thresholds are met.

Design rule for v1 pilots

- Start with one or two workflows, not an entire department.
- Keep the customer system of record in the center of the process; Ruh.ai should orchestrate work around it, not replace it.
- The first goal is trust and measurable delta, not maximum autonomy.

2. Pilot principles and scope discipline

1. Prove one narrow problem at a time. A pilot should target a painful, repetitive workflow with enough volume to measure value, but not so much surface area that change management swamps the result.
2. Read first, write later. Every pilot begins with monitor or draft behavior before any system write-backs are enabled.
3. Approvals are part of the product. Human validation is not a workaround; it is the mechanism that creates trust and allows safe promotion to higher autonomy.

4. Measure against a locked baseline. Time saved, turnaround reduction, approval behavior, and task quality must be measured against the customer's current state rather than rough anecdotes.
5. Treat every pilot as a pack-improvement engine. Every exception, false positive, missing permission, and user complaint should convert into a reusable pack or overlay change.

Pilot archetype	Recommended first workflows	Default overlay profile	Primary buyer	Value proof
Construction Project Operations	Meeting prep, action-item chase, RFI or submittal triage	Procore + SharePoint + Teams	Operations leader, PMO, VDC, project controls	Hours saved per project, overdue actions reduced, faster coordination turns
AI BuildOps	Intake triage, spec drafting, release notes, incident follow-up	GitHub + Linear + Slack + Sentry	CTO, VP Engineering, platform or product ops	Cycle time reduction, fewer manual handoffs, faster incident follow-up

3. Customer qualification and entry criteria

Do not start a design-partner pilot unless the customer can complete the minimum operating loop.

Qualification is a product decision, not only a sales decision. Poor-fit pilots consume implementation effort, distort product priorities, and create false negatives. A customer must demonstrate enough process maturity to let Ruh.ai observe work, retrieve context, receive approvals, and write results back into a trusted system.

Criterion	Why it matters	Minimum threshold for pilot start	Weight
Named executive sponsor	Removes ambiguity when priorities or approvals stall	Required	10
Named workflow owner	Owns current process, approvals, and user feedback	Required	10
System of record in active use	The employee needs a place to read from and write back to	Required	10
Integration access	Admin access, API keys, webhook setup, or app install rights are available	Required	10
Workflow volume	At least one target workflow has enough weekly volume to generate a baseline and a sample set	Required	8
Approval participants	Specific approvers are assigned before write-back behavior is enabled	Required	8
Security and IT readiness	Trust boundary, data handling rules, and allowed environments are approved	Required	10
Baseline metrics available	Current cycle time, manual effort, or backlog levels can be measured	Required	8
Change appetite	Customer is willing to run shadow mode and give weekly feedback	Required	6
Champion user group	At least 2-5 users will use the pilot consistently	Strongly preferred	5

Expansion potential	A second workflow or second team is visible if the pilot succeeds	Preferred	5

Qualification scoring rule

- Required items are binary: if any required item is missing, the pilot does not start.
- Weighted scoring is used to compare otherwise qualified opportunities and prioritize implementation capacity.
- Low-volume or low-discipline customers may still buy later, but they should not shape the v1 pack and overlay roadmap.

4. Pilot operating model and team responsibilities

Every pilot needs two aligned teams: the customer team that owns the workflow and systems of record, and the Ruh.ai team that configures the pack, overlay, policy gates, and success scorecard. The chart below uses a simple RACI model to reduce confusion during the first 90 days.

Work item	Customer sponsor	Workflow owner	IT/admin	Pilot users	Ruh PM	Impl. lead	Integrations	Support/ evals
Pilot charter	A	R	C	I	R	C	I	I
Workflow mapping	I	A/R	C	C	C	R	I	C
Connector setup	I	C	A/R	I	I	C	R	I
Approval matrix	C	A/R	C	I	C	R	I	C
Eval and baseline design	I	C	I	C	C	A	I	R
Weekly scorecard	C	C	I	I	A	R	I	R
Go/no-go gate review	A	R	C	I	R	C	I	C
Incident and rollback response	I	C	C	I	C	A	R	R

Minimum Ruh.ai staffing for a v1 design partner

- One implementation lead who owns timeline, risk tracking, customer coordination, and gate reviews.
- One product manager or product owner who captures exceptions and prioritizes pack and overlay changes.
- Shared integration and support/eval capacity for connector setup, incident triage, and quality measurement.

5. Pilot lifecycle, gates, and deliverables

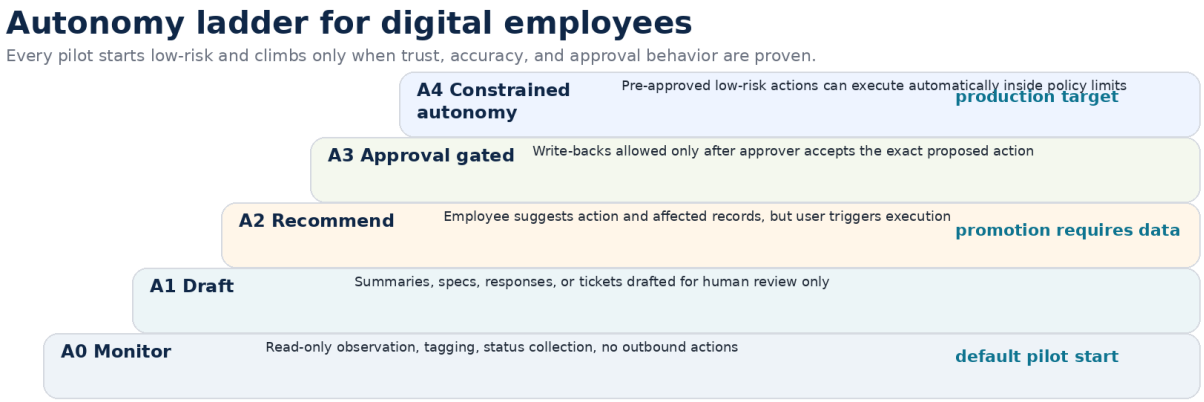

Figure 2. Promotion to higher autonomy is earned through measured behavior, not enabled by default.

Phase 0 - Qualify and charter

Objective: Confirm that the customer, workflow, trust boundary, and baseline data make the pilot worth running.

Key activities: Run readiness assessment, choose the first workflow, define sponsor and workflow owner, and lock the success scorecard.

Exit criteria: Signed pilot charter, target workflow selected, system-of-record access confirmed, and required participants named.

Primary artifacts: Readiness assessment, charter, sponsor map, initial success scorecard.

Gate review questions

- Do we have the right sponsor, owner, and workflow?
- Is there enough volume and access to prove value?
- Would failure teach the roadmap something reusable?

Phase 1 - Discover and baseline

Objective: Capture the live workflow, document the current handoffs, and lock a measurable baseline.

Key activities: Shadow current users, collect example artifacts, classify approval steps, and measure current manual cycle time or effort.

Exit criteria: Workflow map approved, task samples gathered, baseline published, and approval matrix drafted.

Primary artifacts: Workflow map, sample dataset, baseline report, draft approval matrix.

Gate review questions

- Is the current workflow documented with real samples?
- Are the baseline metrics trustworthy?
- Are approval points explicit rather than implied?

Phase 2 - Configure and connect

Objective: Compile the pack and tenant overlay into a customer-specific runtime with connectors, policies, and observability turned on.

Key activities: Provision gateway(s), configure channels, seed memory/context, install connectors, set tool policies, and create initial eval set.

Exit criteria: Dry-run passes, audit logs visible, connector health confirmed, and sandbox or trust-boundary rules verified.

Gate review questions

- Does the runtime respect the trust boundary?
- Can we observe logs, costs, and connector health?
- Is the system safe to enter shadow mode?

Primary artifacts: Compiled overlay, connector auth matrix, policy bundle, eval seed set, dry-run report.	

Phase 3 - Shadow mode **Objective:** Let the digital employee observe or draft work without changing any system of record. **Key activities:** Generate drafts, compare against human output, measure usefulness, collect misses, and tune prompts or pack policies. **Exit criteria:** Enough representative tasks completed, acceptance and quality thresholds met, and no unresolved critical trust issues. **Primary artifacts:** Shadow scorecards, exception log, tuned templates, shadow gate review.	**Gate review questions** • Is draft usefulness high enough to justify write-back risk? • Are misses explainable and fixable? • Do users actually engage with the output?

Phase 4 - Approval-gated write-backs **Objective:** Allow the employee to perform constrained write-backs after explicit human approval. **Key activities:** Enable named approval flows, execute low-risk actions, monitor approval latency, and tune policy restrictions. **Exit criteria:** Three consecutive stable reporting periods, no critical policy failures, and clear value against baseline. **Primary artifacts:** Approval audit trail, action logs, gate review, conversion recommendation.	**Gate review questions** • Are approvals fast enough to preserve workflow value? • Have action errors and policy violations stayed within tolerance? • Is business value clearly above baseline?

Phase 5 - Expand or convert **Objective:** Move a successful pilot into production scope or expand to the next workflow, team, or autonomy level. **Key activities:** Choose expansion lane, finalize support model, package reusable learnings, and agree on production roadmap. **Exit criteria:** Production proposal or scale plan approved, support handoff complete, and pilot closeout report signed off. **Primary artifacts:** Closeout report, production plan, new overlay delta, next-pack backlog.	**Gate review questions** • Which lane creates the best production value: broader users, deeper workflow, or higher autonomy? • Which learnings must become pack defaults? • Can support and ownership shift from pilot mode to production mode?

6. Approval model, action classes, and rollback controls

The pilot should make it impossible to confuse model capability with permission to act. The approval model below is the default for v1 unless a stricter customer policy overrides it. Any move toward constrained autonomy should happen only after the pilot has shown stable quality, user trust, and complete auditability.

Action class	Examples	Default pilot mode	Notes
Read-only observation	Status fetch, summarization, tagging, daily digest	Auto	No state change outside Ruh.ai.
Draft creation	RFI draft, issue spec,	Human accepts draft	Draft can be discarded with

	response template, meeting pack		no side effect.
Low-risk write-back	Create ticket, add comment, update status field, append notes	Named approver required	Must show exact payload before execution.
External communication	Send templated message, notify stakeholder, post summary to channel	Named approver required	Only after channel and audience are pre-approved.
Destructive or high-risk action	Delete, close, archive, contract/cost change, customer-facing commitment	Human-only	Out of scope for v1 pilot autonomy.

- Rollback must be immediate and policy-driven: switch the workflow to monitor-only mode, disable write-back permissions, and keep logs intact for root-cause analysis.
- Every phase gate review should nominate a single operator empowered to stop the pilot or drop the workflow back one autonomy level without debate.
- If a trust or quality incident occurs, the platform team should prefer reducing permissions before changing prompts or models.

7. Workstreams and implementation checklist

Workstream	Owned by	Core tasks	Primary output
Pack and workflow config	Ruh.ai product + implementation	Choose employee roster, enable workflow templates, tune prompts and skill settings	Pilot-specific pack delta
Gateway and trust boundary	Ruh.ai implementation + customer IT	Provision gateway, channel bindings, sandbox mode, network and environment separation	Runtime deployment record
Connectors and credentials	Customer IT + Ruh.ai integrations	Install apps, create tokens, map webhooks, test read and write-back paths	Connector auth matrix
Context and memory	Workflow owner + Ruh.ai product	Seed documents, naming conventions, project/repo mappings, reference datasets	Context seed set
Approvals and policy	Workflow owner + sponsor + Ruh.ai product	Name approvers, define action classes, set escalation and rollback rules	Approval matrix
Evals and scorecard	Ruh.ai support/evals + workflow owner	Create baseline, representative task set, quality rubrics, promotion thresholds	Eval suite + scorecard
Adoption and support	Customer champion + Ruh.ai implementation	User onboarding, feedback capture, weekly office hours, incident channel	User enablement kit

Figure 3. Cadence is part of the pilot design; without regular operating reviews the pack never matures.

8. Scorecard and promotion thresholds

A pilot succeeds only when the scorecard shows that users trust the employee, the workflow improves economically, and the runtime behaves safely and reliably. The scorecard below is the default v1 pattern; each pilot can add pack-specific metrics, but should not remove these four dimensions.

Dimension	Core metrics	Shadow -> approval target	Approval -> expand target	Owner
Business value	Hours saved, task turnaround reduction, overdue backlog reduction	Directional improvement visible	Clear, repeated improvement for 3 reporting periods	Sponsor + workflow owner
Trust and quality	Draft usefulness, factual accuracy, approval acceptance, policy violations	Draft usefulness >= 80%, no critical trust issue	Stable approvals, no critical policy failure	Workflow owner + Ruh.ai product
Technical health	Task success rate, connector uptime, median run latency, audit completeness	No blocking reliability issue	Stable health within agreed operating range	Implementation + support
Adoption	Weekly active users, repeat usage, feedback response rate	Named champions use the system consistently	Usage spreads beyond pilot champions	Customer champion + Ruh.ai PM

Promotion rule of thumb

- Do not promote a workflow from shadow to approval mode on volume alone; it needs representative examples, user trust, and an explicit approver map.
- Do not expand to new teams until the current team uses the workflow repeatedly without implementation hand-holding.
- If business value looks strong but trust is weak, hold the workflow where it is and invest in clarity, retrieval quality, and approval UX.

9. Recommended first-pilot templates

Construction Project Operations pack - recommended first pilot

- Start with one project or one project cluster, not the entire portfolio.
- Choose one or two workflows: meeting pack generation, action-item chase, RFI or submittal triage.

- Use Procore as the system of record, SharePoint for supporting documents, and Teams for user-facing interaction if that is the customer baseline.
- Keep the first value proof simple: reduced coordination time, fewer overdue actions, better draft quality, and less project engineer admin overhead.

Stage	Construction example	Success signal
Shadow	Weekly coordination pack, meeting summary, open-item digest	Project team reads the output before preparing their own version
Approval	Create or update low-risk coordination records after named approval	Team accepts that the system can save real admin time
Expand	Add another workflow or another project team	Learnings are reusable without bespoke logic

AI BuildOps pack - recommended first pilot

- Start with one engineering team or one product area that already uses GitHub, Linear or Jira, Slack, and incident tooling consistently.
- Choose one or two workflows: intake triage, spec drafting, release-note assembly, or post-incident follow-up.
- Use GitHub and the issue tracker as the systems of record; Slack is the collaboration surface, not the source of truth.
- Value proof should focus on shorter coordination loops, cleaner specs, faster follow-up after release, and reduced manual status chasing.

Stage	BuildOps example	Success signal
Shadow	Ticket triage and spec drafts posted for review	PMs and engineers prefer starting from the draft rather than a blank page
Approval	Create or update issues, release notes, or incident follow-ups after human approval	The team trusts the orchestrated handoff and uses it repeatedly
Expand	Add builder orchestration, eval generation, or another team	The operating model scales beyond one lead user

10. Common failure modes, mitigations, and closeout

Failure mode	Why it happens	Mitigation or closeout rule
Pilot scope is too broad	Customer tries to include too many workflows or user groups before trust is earned	Reduce to one workflow and re-baseline before continuing
No reliable owner	Sponsor is present but workflow owner cannot make day-to-day decisions	Pause or stop; the pilot cannot be run by committee
Low data discipline	Documents, tickets, or project records are too inconsistent to support retrieval or write-back	Add normalization work or choose a simpler workflow
Approval bottlenecks	Named approvers do not respond quickly enough to keep the workflow useful	Narrow action class, add backup approver, or stay in draft mode
Trust incident	Output contains a material factual error or policy violation	Drop to monitor mode, preserve evidence, run root-cause review, update pack/policy before restart
No measurable delta	The workflow runs but value against baseline is unclear	Tighten scorecard or end the pilot cleanly rather than pretending success

- A clean pilot closeout is acceptable. If the workflow is low-value, the customer is not ready, or the trust boundary is wrong, end the pilot, summarize the learnings, and feed the roadmap.
- The closeout packet should always include: scorecard summary, unresolved risks, pack and overlay deltas, support handoff recommendation, and the commercial conversion decision.

11. Standard pilot artifact set

Artifact	Purpose
Pilot charter	Defines scope, sponsor, owners, success scorecard, and decision cadence.
Readiness assessment	Confirms trust boundary, system access, workflow volume, and qualification status.
Workflow map and baseline	Documents the current state and the numbers the pilot must improve.
Connector auth matrix	Tracks each connected system, owner, environment, credential method, and expiry.
Approval matrix	Defines action classes, approvers, escalation, and rollback authority.
Eval seed set and scorecard	Measures quality, trust, technical health, and adoption across phases.
Weekly operating review	Captures decisions, incidents, requested changes, and go/no-go status.
Pilot closeout report	Explains what worked, what failed, what changed in the pack/overlay, and whether to expand.

Appendix A - Go / no-go checklist

Check	Status at kickoff
Sponsor, workflow owner, and IT/admin contact are confirmed.	Pending / Pass / Blocked
Target workflow has baseline metrics and representative samples.	Pending / Pass / Blocked
Runtime trust boundary and environment are approved.	Pending / Pass / Blocked
Connectors needed for the pilot are configured and tested.	Pending / Pass / Blocked
Approval classes and named approvers are documented.	Pending / Pass / Blocked
Shadow-mode scorecard and promotion thresholds are agreed.	Pending / Pass / Blocked
Rollback owner and incident channel are defined.	Pending / Pass / Blocked